Ananya Joshi

Jaipur, Rajasthan | (+91) 8003944392 | aj.ananyajoshi@gmail.com | linkedin.com/in/ananyajoshi | github.com/Ananya-d14

Data Scientist and analyst with 3+ years of experience in LLM evaluation, structured reasoning, and technical content creation. Built and benchmarked LLM-based systems with hand-labelled evaluation sets, citation-faithfulness verification, and LLM-judge calibration against human ratings. Strong analytical and writing skills, with experience producing reproducible technical content for 200+ learners and validating model outputs against source evidence.

CORE STRENGTHS

Analytical & Research: Structured problem decomposition, claim validation against sources, logical reasoning, quantitative analysis, basic arithmetic/data interpretation
LLM Work: RAG system design, prompt engineering, LLM-judge evaluation, benchmark construction, citation faithfulness checks, hallucination detection, fine-tuning, sentiment classification
Writing & Communication: Technical documentation, instructional content for bootcamp learners, structured annotation, stakeholder reporting, English (professional working proficiency MSc in Ireland)
Tools: Python, SQL, Excel, Google Suite, Hugging Face, LangGraph, Tableau, Jupyter, Git, FastAPI

EXPERIENCE

Skill Arbitrage **Gurugram, India**
Data Science Consultant *Oct 2024 – Present*

- Authored reusable Python data-analysis modules and instructional content adopted across 10+ bootcamp modules by 200+ learners, with detailed step-by-step explanations and annotated examples.
- Designed and trained a custom language model for customer-email sentiment classification, reaching 75–78% accuracy; iteratively refined the labelling rubric to reduce ambiguous edge cases.
- Built and led a 7-person freelance team delivering 5 client projects in 6 months wrote project briefs, structured deliverable formats, and reviewed analytical outputs for correctness and clarity.

Solkraft Technology **Jaipur, India**
Data Analyst *Sep 2023 – Oct 2024*

- Validated terabytes of solar-energy data against expected patterns; designed automated checks that improved data accuracy by 20% and surfaced anomalies for human review.
- Translated business questions from non-technical stakeholders into structured SQL analyses and dashboards, contributing to 10% efficiency gains and 7% cost reduction.

Housing and Urban Development Corporation Ltd (HUDCO) **New Delhi, India**
Data Scientist & Analyst Intern *Jan 2023 – Aug 2023*

- Wrote 30+ structured weekly stakeholder reports during a large ERP implementation, ensuring 100% data consistency across analytical narratives shared with leadership.
- Built an ML-powered recommendation engine improving conversion by 30%, and produced clear written documentation explaining model logic for non-technical decision-makers.

SELECTED PROJECTS

Financial Filings Agent LLM evaluation & multi-hop reasoning over SEC filings | *Python, LangGraph, LLM-as-judge, benchmarking*

- Designed and hand-labelled a 100-question benchmark across 4 difficulty tiers (single-fact lookup → multi-company multi-hop reasoning) to systematically evaluate LLM accuracy on financial-document tasks.
- Built an LLM-as-judge evaluation harness, calibrated against human ratings using Cohen's κ, and ran a 5-config ablation tracking accuracy, citation faithfulness, latency, and cost exactly the kind of structured reasoning work that catches model failure modes.
- Engineered a citation-verification step: every numeric claim the LLM produces is programmatically checked against source text before output, eliminating an entire class of hallucinations.

Credit Risk Scoring API | *Analytical reasoning, business-outcome modelling*

- Reframed loan-approval as a business-outcome problem: instead of optimising accuracy, tuned the decision threshold against a profit/loss matrix (+$100 per good loan, −$500 per default) surfacing the cost of each modelling assumption.
- Wrote 6 reproducible notebooks documenting the full reasoning chain feature selection rationale, model comparisons, SHAP explanations for analysts, and a FICO-style scorecard for non-technical readers.

Sentiment Classification on IMDB Reviews | *LSTM, model benchmarking*

- Benchmarked three model architectures (Dense, CNN, LSTM) on the same 50K-review split to quantify exactly how much sequence modelling buys for the task LSTM reached 85.8% test accuracy vs Dense at 74.5%.

EDUCATION **MSc. Data Science and Analytics, University College Cork** Cork, Ireland | Sep 2021 – Oct 2022